

BARTHOLOMEW AI

Autonomous CI/CD Failure Auto-Fix & Mechanical Verification Engine

Platform: bartholomew.info
Contact: help@bartholomew.info
Version: Executive Summary v2.0

01. Personal Story & Founder Background

I am a systems builder driven by an obsession with operational reliability. My background is rooted in building resilient, mission-critical infrastructure rather than academic pedigree or corporate theater. Throughout years of hands-on software development, I experienced the universal friction of enterprise engineering: brilliant developers spending 20% of their workweeks trapped in "CI/CD hell"—debugging broken builds, fixing upstream dependency deprecations, and manually triaging flaky tests at 2:00 AM.

When the generative AI boom arrived, the market flooded with shallow chatbot wrappers that promised to "write all your code." But in practice, ungrounded Large Language Models introduce subtle bugs, hallucinate non-existent libraries, and pollute repositories with untested changes. I set out to build the exact opposite: not an AI chatbot that guesses code, but an **autonomous mechanical verification engine** that mathematically proves a fix works before a human engineer ever opens a pull request.

02. The Consensus View We Absolutely Agree With

"Computers execute instructions with absolute literalism, never with good intentions—and deterministic verification is the only ground truth in software engineering."

While modern AI hype argues that probabilistic neural networks will replace traditional software mechanics, we firmly agree with the foundational computer science consensus: **probability is never a substitute for proof.**

An AI model guessing a patch with 90% confidence leaves a 10% risk of critical production outages. Bartholomew is built on the belief that traditional compiler AST parsing, isolated container sandboxing, binary test pass/fail exit codes, and cryptographic signatures are timeless. The future of autonomous coding belongs to deterministic mechanical harnesses that constrain and verify AI outputs with zero tolerance for error.

03. How Bartholomew Works (Layman's Guide for Investors)

Imagine your company's software is a high-speed Formula 1 race car. Every time an engineer changes a part, it goes through a routine safety inspection (CI/CD tests). If something breaks, the car stops, the entire pit crew freezes, and engineers spend hours looking for the broken bolt.

1,000,000

Verified Test Cycles

100.00%

Zero Regression Pass

1.14 μ s

Kernel Trajectory Intercept

3 Lines

Average Surgical Delta

Bartholomew acts as an Autonomous Robotic Pit Crew in 4 deterministic steps:

1. **The Isolated Test Track (Ingest & Reproduce):** The second a test fails in GitHub, Bartholomew clones the failure into a private, isolated digital sandbox and recreates the exact crash so no human has to guess what went wrong.
2. **The Laser Scalpel (Compiler AST Surgery):** Instead of letting AI blindly rewrite whole files, Bartholomew examines the code's grammatical tree (Abstract Syntax Tree) and calculates the *exact minimal 3-line surgical fix* required.
3. **The 100-Point Crash Test (Pre-Flight Verification):** Bartholomew runs the entire battery of company tests against the fix inside the sandbox. **If 100% of tests do not pass green, the fix is immediately rejected.**
4. **The Cryptographic Wax Seal (Ed25519 Receipt):** Once verified, Bartholomew signs the execution proof with tamper-proof mathematical cryptography (RFC 8785 Ed25519) and submits a completed, green Pull Request ready to merge.

04. The "1-of-1" Moat: Addressing Naysayers & Skeptics

When critics dismiss automated tools as "just an AI script that can be cloned in a weekend," they fail to understand the difference between a toy prompt and a mechanical verification engine.

✗ The Skeptic's Assumption

"Anyone can copy-paste a stack trace into ChatGPT and ask it to fix a bug in 20 lines of Python."

✓ The Bartholomew Reality

Raw LLM prompts hallucinate imports, alter unrelated code, and fail tests. Bartholomew executes isolated sandbox reproduction proofs, compiler AST analysis, and 100% pre-flight test suites.

✗ The Skeptic's Assumption

"AI coding tools are too slow and ungrounded for enterprise security policies."

✓ The Bartholomew Reality

Our compiled Go daemon intercepts trajectories in **1.14 microseconds** (11.98M ops/sec), enforcing real-time OWASP Top-10 kill switches before dangerous commands execute.

Empirical Defense: We subjected Bartholomew to a rigorous **1,000,000-cycle continuous stress benchmark** across 12 parallel CPU cores: **0 failures (0.00000%)**, **100.0000% reliability**, and **28,880 cryptographic ops/sec**. This is compiled systems engineering, not a fragile prototype.

05. Market Opportunity & Big Tech Integration

At companies like Google, Microsoft, Meta, and Amazon, an engineering organization of 500 developers loses over **\$20,000,000 annually** in developer payroll to broken CI builds and dependency triage. Bartholomew targets this massive enterprise inefficiency:

Target Partner	Integration Opportunity	Core Value Proposition
Google Cloud / OSS	Cloud Build native add-on & Open-Source Monorepo Auto-Remediator	Auto-resolves Python 3.13/3.14 AST deprecations across repos like <i>Google Python Fire</i> & <i>TensorFlow</i> .
Microsoft / GitHub	GitHub Actions Marketplace "Auto-Resolve" Engine	Turns broken CI runs into instant green PRs, increasing developer velocity and GitHub compute usage.
OpenAI / Anthropic	Zero-latency Trajectory Firewall for Autonomous Agents (Operator)	Provides the cryptographic Ed25519 audit receipts and POSIX kill-switch Fortune 500 CISOs require.

06. Implementation Roadmap & Budget Allocation

- **Phase 1 (Months 1–3): Open-Source Ecosystem Ingestion:** Deploy automated zero-regression PR bots across high-profile open-source repositories to build brand authority and prove zero-defect accuracy.
- **Phase 2 (Months 4–6): GitHub Marketplace & Self-Serve SaaS:** Launch self-serve GitHub App (\$49/repo/month) with automated webhook ingestion and CLI tool.
- **Phase 3 (Months 7–12): Enterprise Private Cloud:** Deploy dedicated VPC sandbox runners and Ed25519 audit compliance suites for SOC2/ISO27001 enterprise customers (\$499–\$2,500/month).

Expense Category	Description	Target Allocation
Sandboxed Cloud Runners	High-concurrency hermetic container clusters for multi-repo test execution	\$30,000
Security & Formal Audits	Independent cryptographic attestation & penetration testing	\$25,000
Developer Pilot Programs	Enterprise onboarding, developer documentation, and ecosystem grants	\$25,000
Operational Reserve	Edge CDN hosting, domain infrastructure, and contingency runway	\$20,000
Total Grant / Seed Request	12-Month Runway for Full-Time Deployment	\$100,000